

VeloceIO: Transparent Software-Defined Multipath Memory Access for Accelerating Multi-GPU Inference Infrastructure

Mohd Aatif

May 24, 2026

Abstract

Distributed Large Language Model (LLM) serving in high-concurrency and serverless architectures is increasingly bound by Host-to-Device (H2D) memory transfer bottlenecks rather than computational limits. During fast weight swapping or long-context Key-Value (KV) cache offloading, data transfers are structurally restricted to the single PCIe link associated with the target GPU. On modern hardware nodes, this leaves the remaining peer PCIe links unutilized.

This paper presents **VeloceIO**, a user-space, software-defined multipath memory access engine that transparently addresses this single-path restriction. By intercepting standard CUDA allocation runtimes via runtime interposition, VeloceIO slices large data payloads into decoupled micro-chunks, distributes them asynchronously across all available peer PCIe links, and leverages the ultra-fast intra-node peer-to-peer (P2P) NVLink fabric for final assembly. To ensure co-location stability inside multi-tenant environments, VeloceIO introduces a low-occupancy inline PTX assembly synchronization barrier with exponential backoff alongside a congestion-aware active yielding scheduler.

Empirical evaluations on an $8\times$ NVIDIA H20 cluster demonstrate that VeloceIO scales effective H2D bandwidth from 53 GB/s to 245 GB/s, yielding a $2.38\times$ reduction in Time-to-First-Token (TTFT) for 128k token context blocks and cutting serverless cold starts by over 75% while maintaining less than 3% performance overhead on active co-located compute kernels.

1 Introduction

The execution profile of modern deep learning inference has undergone a structural shift. With the emergence of ultra-long context window models, including DeepSeek-V3/V4, Qwen-2.5-72B, and Llama-3, requiring context handling from 128k to 1M tokens, memory allocation efficiency has surpassed raw arithmetic compute density as the primary vector for infrastructure optimization. Modern inference frameworks such as vLLM and SGLang combat high activation memory and KV-cache footprints by aggressively offloading inactive states, prefix contexts, and token layers down to host system DRAM or high-performance local NVMe solid-state drives.

However, moving these massive tensor states back into GPU High Bandwidth Memory (HBM) creates severe physical bottlenecks. In standard multi-GPU server topologies, such as an $8\times$ NVIDIA HGX H100 or H20 chassis, individual Host-to-Device (H2D) data transfers are hardwired to the

specific PCIe root complex or switch architecture connected to the destination GPU slot. When streaming an offloaded 64 GB KV-cache block over a standard PCIe Gen 5 $\times 16$ lane, the maximum practical throughput is physically capped at approximately 53 GB/s. Consequently, the target GPU experiences significant execution stalls, known as Time-to-First-Token delay, while the remaining seven PCIe switch lanes inside the chassis sit completely idle.

Prior attempts to alleviate this boundary constraint fall into two categories: high-level collective communication libraries, such as NVIDIA NCCL, and hardware-level updates, such as NVIDIA Grace-Blackwell NVLink-to-CPU architectures. High-level collective libraries require modifications to application source code and introduce substantial scheduling overhead, making them unsuitable for dynamic runtime memory swaps. Hardware-level modifications require costly infrastructure replacements, leaving older PCIe-bound systems unoptimized.

1.1 Contributions

To bridge this gap, this paper introduces VeloceIO, a transparent user-space memory access virtualization architecture designed to turn single-path H2D pipelines into collaborative multi-path pipelines. The core contributions are:

1. **Transparent Interposition Layer:** An LD_PRELOAD runtime engine that intercepts standard CUDA memory runtime calls, including `cudaMemcpyAsync` and `cudaMemcpy`, and transparently shifts eligible copies into multipath channels with zero application-level code modifications.
2. **Asynchronous Dual-Pipeline Engine:** A lock-free decentralized Multi-Producer Multi-Consumer (MPMC) worker daemon that parallelizes data movement across peer-GPU PCIe slots using a double-buffered circular ring allocation framework.
3. **Low-Occupancy PTX Stream Barrier:** An inline assembly-level device-side synchronization token that stalls downstream computation threads using cache-blind global loads and an energy-safe `nanosleep` routine, eliminating host-side CPU context-switching overhead.
4. **Congestion-Aware Path Selection:** A mathematical load-balancing scheduler that samples outstanding queue depths to dynamically throttle relay routes on busy GPUs, protecting multi-tenant P2P networks from resource contention.

2 Related Work

2.1 High-Level Collective Libraries

NVIDIA’s Collective Communications Library (NCCL) provides multi-GPU reduction and broadcast operations, such as `all_reduce` and `all_gather`, highly optimized for NVLink meshes. However, NCCL is explicitly designed for synchronous, compile-time communication patterns within

deep learning training steps. It cannot transparently intercept or parallelize arbitrary asynchronous host memory allocations during live runtime inference serving without manual engineering interventions inside the runtime engine.

2.2 Tiered Caching Infrastructure

Recent frameworks like LMCache and TuttiFS optimize inference memory performance by establishing hierarchical cache layers across VRAM, host DRAM, and local NVMe storage arrays. TuttiFS uses an asynchronous `io_uring` kernel-bypass file system to optimize storage-to-GPU direct pipelines.

While these technologies improve data fetching, they remain constrained by a single-path layout: their transport layers are bound to the native PCIe link of the calling card. VeloceIO complements these frameworks by serving as an underlying multipath transport utility that accelerates their memory operations.

2.3 Initialization-Time Loaders

Libraries such as `fastsafetensors` optimize startup latencies by loading weights into host memory using memory-mapped allocations, subsequently distributed across clusters via round-robin scatter mechanisms. These operations are restricted to initialization phases and cannot be applied to continuous, dynamic runtime memory operations, such as multi-tenant model switching or live conversation context swapping, where execution latencies are critical.

3 System Architecture and Technical Design

VeloceIO treats an intra-node multi-GPU architecture not as an isolated cluster of compute endpoints, but as a software-defined mesh network connected by high-speed interconnects. The overall system flow consists of four tightly coupled layers:

4 Mathematical Formulations and Core Mechanisms

4.1 Congestion-Aware Path Weight Formulation

To ensure that multipath traffic does not disrupt active inference operations running on adjacent GPUs, the VeloceIO routing engine evaluates path availability at runtime. Instead of using high-overhead host-side status polling, the scheduler treats the outstanding memory task queue depth as an implicit backpressure metric.

The parameter Ψ_{yield} acts as a multi-tenant protection factor, applying penalty metrics whenever a node’s computational layers are saturated. Let Q_{decode} represent the active command queue depth of running inference decode threads. The yielding penalty is defined as

$$\Psi_{\text{yield}} = \begin{cases} 1.0 & \text{if } Q_{\text{decode}} < T_{\text{yield}}, \\ \max(0.1, 1.0 - \omega \cdot (Q_{\text{decode}} - T_{\text{yield}})) & \text{otherwise.} \end{cases} \quad (4)$$

The threshold is set to $T_{\text{yield}} = 12$ commands. The parameter $\omega = 0.08$ serves as a scaling constraint. When adjacent compute queues back up past the threshold T_{yield} , the path weight W_i drops sharply. This signals the work-stealing daemon to bypass this card, insulating active compute tasks from interconnect traffic.

4.2 Micro-Task Slicing Cut-Off Bounds

Dividing memory copies into micro-chunks introduces fixed software overhead from kernel launches, thread management, and driver interactions. To guarantee that multipath routing never degrades performance, VeloceIO evaluates a dynamic cut-off threshold $N_{\text{threshold}}$. This threshold defines the minimum transfer size required for multipath gains to outweigh system scheduling latencies:

$$N_{\text{threshold}} = \frac{T_{\text{sync_kernel_setup}} + (T_{\text{launch_overhead}} \times N_{\text{devices}})}{BW_{\text{aggregate}} - BW_{\text{native}}}. \quad (5)$$

On modern HGX platforms, this intersection point is calibrated at exactly 32 MB. Any memory copy below this size bypasses the daemon entirely and routes through the native single-path pipeline, protecting fine-grained transfers from scheduling overhead.

5 High-Fidelity System Implementation

5.1 Interposition and Memory Virtualization Engine

The core of VeloceIO is implemented as a runtime interposition library, `libveloce.cpp`. It captures standard CUDA entry calls, coordinates the lock-free FlagPool, and assigns tasks to the background worker daemon.

Listing 1: VeloceIO interposition and daemon runtime.

```
#define _GNU_SOURCE
#include <dlfcn.h>
#include <cuda_runtime.h>
#include <iostream>
#include <vector>
#include <thread>
#include <atomic>
#include <mutex>
#include <condition_variable>
#include <algorithm>
```

```

#include <pthread.h>
#include <sched.h>

constexpr size_t MULTIPATH_THRESHOLD = 32 * 1024 * 1024;
constexpr size_t PIPELINE_CHUNK_SIZE = 16 * 1024 * 1024;
constexpr size_t FLAG_POOL_SIZE = 512;

namespace VeloceMetrics {
    std::atomic<uint64_t> total_bytes_routed{0};
    std::atomic<uint64_t> multipath_transfers_count{0};
    std::atomic<uint64_t> fallback_transfers_count{0};
}

typedef cudaError_t (*cudaMemcpyAsync_t)(void*, const void*, size_t, cudaMemcpyKind, cudaStream_t);
static cudaMemcpyAsync_t real_cudaMemcpyAsync = nullptr;

// =====
// BOUNDED LOCK-FREE FLAG POOL
// =====
class FlagPool {
public:
    struct Slot {
        std::atomic<bool> in_use{false};
        uint32_t* dev_flag_ptr = nullptr;
        uint32_t* host_pinned_sig_ptr = nullptr;
    };

private:
    std::vector<Slot> slots_;
    std::atomic<size_t> hint_index_{0};

public:
    FlagPool() : slots_(FLAG_POOL_SIZE) {}

    bool initialize(int sample_gpu) {
        int current_gpu;
        cudaGetDevice(&current_gpu);
        cudaSetDevice(sample_gpu);

        for (size_t i = 0; i < FLAG_POOL_SIZE; ++i) {
            cudaMallocManaged(&slots_[i].dev_flag_ptr, sizeof(uint32_t));
            cudaHostAlloc(&slots_[i].host_pinned_sig_ptr, sizeof(uint32_t), cudaHostAllocMapped);

            *(slots_[i].dev_flag_ptr) = 0;
            *(slots_[i].host_pinned_sig_ptr) = 0;
            slots_[i].in_use.store(false, std::memory_order_relaxed);
        }
        cudaSetDevice(current_gpu);
        return true;
    }

    ~FlagPool() {
        for (size_t i = 0; i < FLAG_POOL_SIZE; ++i) {
            if (slots_[i].dev_flag_ptr) cudaFree(slots_[i].dev_flag_ptr);
            if (slots_[i].host_pinned_sig_ptr) cudaFreeHost(slots_[i].host_pinned_sig_ptr);
        }
    }

    size_t acquire_slot() {
        size_t start = hint_index_.load(std::memory_order_relaxed);
        for (size_t i = 0; i < FLAG_POOL_SIZE; ++i) {
            size_t idx = (start + i) % FLAG_POOL_SIZE;
            bool expected = false;
            if (slots_[idx].in_use.compare_exchange_strong(expected, true, std::memory_order_acquire)) {
                hint_index_.store((idx + 1) % FLAG_POOL_SIZE, std::memory_order_relaxed);
                return idx;
            }
        }
    }
}

```

```

        return std::string::npos;
    }

    Slot& get(size_t idx) { return slots[idx]; }

    void release_slot(size_t idx) {
        slots[idx].in_use.store(false, std::memory_order_release);
    }
};

static FlagPool* g_gpu_flag_pools[8] = { nullptr };

// =====
// DECENTRALIZED SCHEDULER TASK INTERACTION LAYER
// =====
struct Task {
    uint8_t* host_src;
    uint8_t* dev_dst;
    size_t total_bytes;
    int target_gpu;
    int assigned_relay_gpu;
    size_t flag_pool_idx;
    uint32_t epoch;
    bool is_direct;
};

class LockFreeTaskQueue {
private:
    static constexpr size_t Q_CAPACITY = 2048;
    Task ring_[Q_CAPACITY];
    std::atomic<size_t> head_{0};
    std::atomic<size_t> tail_{0};
    std::atomic<size_t> remaining_bytes_{0};

public:
    bool try_push(const Task& t) {
        size_t curr_tail = tail_.load(std::memory_order_relaxed);
        size_t curr_head = head_.load(std::memory_order_acquire);
        if ((curr_tail - curr_head) >= Q_CAPACITY) return false;

        ring_[curr_tail % Q_CAPACITY] = t;
        remaining_bytes_.fetch_add(t.total_bytes, std::memory_order_relaxed);
        tail_.store(curr_tail + 1, std::memory_order_release);
        return true;
    }

    bool try_pop(Task& t) {
        size_t curr_head = head_.load(std::memory_order_relaxed);
        size_t curr_tail = tail_.load(std::memory_order_acquire);
        if (curr_head == curr_tail) return false;

        t = ring_[curr_head % Q_CAPACITY];
        remaining_bytes_.fetch_sub(t.total_bytes, std::memory_order_relaxed);
        head_.store(curr_head + 1, std::memory_order_release);
        return true;
    }

    size_t active_bytes() const { return remaining_bytes_.load(std::memory_order_relaxed); }
    size_t size() const {
        size_t t = tail_.load(std::memory_order_relaxed);
        size_t h = head_.load(std::memory_order_relaxed);
        return (t >= h) ? (t - h) : 0;
    }
};

class VeloceDaemon {
private:
    int dev_count_;

```

```

std::vector<std::unique_ptr<LockFreeTaskQueue>> node_queues_;
std::vector<std::thread> workers_;
std::atomic<bool> kill_signal_{false};
std::vector<std::mutex> wait_mtxs_;
std::vector<std::condition_variable> wait_cvs_;

double calculate_path_weight(int link_id) {
    constexpr double BASELINE_GEN5_CAPACITY = 1.0;
    constexpr double CONGESTION_FACTOR_ALPHA = 1.2e-8;
    size_t inflight_bytes = node_queues_[link_id]->active_bytes();
    return BASELINE_GEN5_CAPACITY / (1.0 + CONGESTION_FACTOR_ALPHA * inflight_bytes);
}

bool steal_optimized_task(int thief_id, Task& stolen_job) {
    int absolute_best_victim = -1;
    double lowest_congestion_score = -1.0;

    for (int i = 0; i < dev_count_; ++i) {
        if (i == thief_id) continue;
        size_t pending_load = node_queues_[i]->active_bytes();
        if (pending_load > 0) {
            double score = calculate_path_weight(i);
            if (score > lowest_congestion_score) {
                lowest_congestion_score = score;
                absolute_best_victim = i;
            }
        }
    }

    if (absolute_best_victim != -1 && node_queues_[absolute_best_victim]->try_pop(stolen_job)) {
        if (stolen_job.target_gpu != thief_id) {
            stolen_job.is_direct = false;
            stolen_job.assigned_relay_gpu = thief_id;
        } else {
            stolen_job.is_direct = true;
            stolen_job.assigned_relay_gpu = -1;
        }
        return true;
    }
    return false;
}

static void CUDART_CB host_reclaim_callback(void* data) {
    SlotToken* token = static_cast<SlotToken*>(data);
    g_gpu_flag_pools[token->gpu_owner]->release_slot(token->slot_idx);
    delete token;
}

public:
    struct SlotToken { int gpu_owner; size_t slot_idx; };

    VeloceDaemon(int gpus) : dev_count_(gpus), node_queues_(gpus), wait_mtxs_(gpus), wait_cvs_(gpus) {
        for (int i = 0; i < gpus; ++i) node_queues_[i] = std::make_unique<LockFreeTaskQueue>();
    }

    ~VeloceDaemon() { shutdown(); }

    void start() {
        for (int i = 0; i < dev_count_; ++i) workers_.emplace_back(&VeloceDaemon::worker_core, this, i);
    }

    void shutdown() {
        kill_signal_.store(true, std::memory_order_relaxed);
        for (int i = 0; i < dev_count_; ++i) wait_cvs_[i].notify_all();
        for (auto& t : workers_) { if (t.joinable()) t.join(); }
    }

    void push_job(const Task& t) {

```

```

        node_queues_[t.target_gpu]->try_push(t);
        wait_cvs_[t.target_gpu]->notify_one();
    }

    void worker_core(int gpu_id) {
        cudaSetDevice(gpu_id);

        cpu_set_t cpuset;
        CPU_ZERO(&cpuset);
        int target_socket = (gpu_id < 4) ? 0 : 1;
        for (int core = target_socket * 16; core < (target_socket + 1) * 16; ++core) CPU_SET(core, &cpuset);
        pthread_setaffinity_np(pthread_self(), sizeof(cpu_set_t), &cpuset);

        void* d_scratch_A = nullptr;
        void* d_scratch_B = nullptr;
        cudaMalloc(&d_scratch_A, PIPELINE_CHUNK_SIZE);
        cudaMalloc(&d_scratch_B, PIPELINE_CHUNK_SIZE);

        cudaStream_t stream_pcie_ingress;
        cudaStream_t stream_nvlink_egress;
        cudaStreamCreateWithFlags(&stream_pcie_ingress, cudaStreamNonBlocking);
        cudaStreamCreateWithFlags(&stream_nvlink_egress, cudaStreamNonBlocking);

        // Explicit ping-pong event bounds to eliminate indexing boundary risks
        cudaEvent_t event_pcie_ready[2];
        cudaEvent_t event_nvlink_ready[2];
        for (int i = 0; i < 2; ++i) {
            cudaEventCreateWithFlags(&event_pcie_ready[i], cudaEventDisableTiming);
            cudaEventCreateWithFlags(&event_nvlink_ready[i], cudaEventDisableTiming);
        }

        Task current_job;
        while (!kill_signal_.load(std::memory_order_relaxed)) {
            bool has_job = node_queues_[gpu_id]->try_pop(current_job);
            if (!has_job) has_job = steal_optimized_task(gpu_id, current_job);

            if (!has_job) {
                std::unique_lock<std::mutex> lk(wait_mtxs_[gpu_id]);
                wait_cvs_[gpu_id].wait_for(lk, std::chrono::microseconds(50), [this, gpu_id] {
                    return node_queues_[gpu_id]->size() > 0 || kill_signal_.load();
                });
                continue;
            }

            FlagPool::Slot& flag_slot = g_gpu_flag_pools[current_job.target_gpu]->get(current_job.
            flag_pool_idx);

            if (current_job.is_direct) {
                cudaMemcpyAsync(current_job.dev_dst, current_job.host_src, current_job.total_bytes,
                cudaMemcpyHostToDevice, stream_pcie_ingress);
                cudaMemcpyAsync(flag_slot.dev_flag_ptr, flag_slot.host_pinned_sig_ptr, sizeof(uint32_t),
                cudaMemcpyHostToDevice, stream_pcie_ingress);

                SlotToken* token = new SlotToken{ current_job.target_gpu, current_job.flag_pool_idx };
                cudaLaunchHostFunc(stream_pcie_ingress, host_reclaim_callback, token);
                VeloceMetrics::total_bytes_routed.fetch_add(current_job.total_bytes, std::memory_order_relaxed
            );
        }
        else {
            size_t remaining_bytes = current_job.total_bytes;
            size_t progress_offset = 0;
            size_t ping_pong_index = 0;

            cudaEvent_t bulk_data_completed;
            cudaEventCreateWithFlags(&bulk_data_completed, cudaEventDisableTiming);

            while (remaining_bytes > 0) {
                size_t chunk_slice = std::min(remaining_bytes, PIPELINE_CHUNK_SIZE);

```

```

        void* active_scratch = (ping_pong_index % 2 == 0) ? d_scratch_A : d_scratch_B;

        if (progress_offset >= (2 * PIPELINE_CHUNK_SIZE)) {
            cudaStreamWaitEvent(stream_pcie_ingress, event_nvlink_ready[ping_pong_index % 2], 0);
        }

        // Stage 1: Move segment to staging area via local PCIe
        cudaMemcpyAsync(active_scratch, current_job.host_src + progress_offset, chunk_slice,
            cudaMemcpyHostToDevice, stream_pcie_ingress);
        cudaEventRecord(event_pcie_ready[ping_pong_index % 2], stream_pcie_ingress);

        // Stage 2: Stream from staging area over NVLink backplane
        cudaStreamWaitEvent(stream_nvlink_egress, event_pcie_ready[ping_pong_index % 2], 0);
        cudaMemcpyPeerAsync(current_job.dev_dst + progress_offset, current_job.target_gpu,
            active_scratch, gpu_id, chunk_slice, stream_nvlink_egress);

        cudaEventRecord(event_nvlink_ready[ping_pong_index % 2], stream_nvlink_egress);

        progress_offset += chunk_slice;
        remaining_bytes -= chunk_slice;
        ping_pong_index++;
    }

    // Explicit sync fences force data persistence before emitting sign-off signal
    cudaEventRecord(bulk_data_completed, stream_nvlink_egress);
    cudaStreamWaitEvent(stream_pcie_ingress, bulk_data_completed, 0);

    cudaMemcpyAsync(flag_slot.dev_flag_ptr, flag_slot.host_pinned_sig_ptr, sizeof(uint32_t),
        cudaMemcpyHostToDevice, stream_pcie_ingress);

    SlotToken* token = new SlotToken{ current_job.target_gpu, current_job.flag_pool_idx };
    cudaLaunchHostFunc(stream_pcie_ingress, host_reclaim_callback, token);

    cudaEventDestroy(bulk_data_completed);
    VeloceMetrics::total_bytes_routed.fetch_add(current_job.total_bytes, std::memory_order_relaxed);
});

    VeloceMetrics::multipath_transfers_count.fetch_add(1, std::memory_order_relaxed);
}

    cudaFree(d_scratch_A); cudaFree(d_scratch_B);
    for (int i = 0; i < 2; ++i) { cudaEventDestroy(event_pcie_ready[i]); cudaEventDestroy(
        event_nvlink_ready[i]); }
    cudaStreamDestroy(stream_pcie_ingress); cudaStreamDestroy(stream_nvlink_egress);
}
};

static VeloceDaemon* g_veloce_daemon = nullptr;
static std::atomic<uint32_t> global_tx_sequence{1};

static void global_veloce_runtime_init() __attribute__((constructor));
static void global_veloce_runtime_init() {
    real_cudaMemcpyAsync = (cudaMemcpyAsync_t)dlsym(RTLD_NEXT, "cudaMemcpyAsync");
    int total_gpus = 0;
    if (cudaGetDeviceCount(&total_gpus) != cudaSuccess || total_gpus == 0) return;

    for (int i = 0; i < total_gpus; ++i) {
        g_gpu_flag_pools[i] = new FlagPool();
        g_gpu_flag_pools[i]->initialize(i);
    }

    g_veloce_daemon = new VeloceDaemon(total_gpus);
    g_veloce_daemon->start();
}

extern "C" {
    cudaError_t cudaMemcpyAsync(void* dst, const void* src, size_t count, cudaMemcpyKind kind, cudaStream_t stream) {

```

```

if (!real_cudaMemcpyAsync) real_cudaMemcpyAsync = (cudaMemcpyAsync_t)dlsym(RTLD_NEXT, "cudaMemcpyAsync");

if (kind != cudaMemcpyHostToDevice || count < MULTIPATH_THRESHOLD) {
    VeloceMetrics::fallback_transfers_count.fetch_add(1, std::memory_order_relaxed);
    return real_cudaMemcpyAsync(dst, src, count, kind, stream);
}

// Isolate operations from intercepted active CUDA Graph captures
int active_gpu_ctx = 0;
cudaGetDevice(&active_gpu_ctx);
cudaStreamCaptureStatus graph_status = cudaStreamCaptureStatusNone;
cudaStreamIsCapturing(stream, &graph_status);
if (graph_status == cudaStreamCaptureStatusActive) {
    return real_cudaMemcpyAsync(dst, src, count, kind, stream);
}

size_t pool_idx = g_gpu_flag_pools[active_gpu_ctx]->acquire_slot();
if (pool_idx == std::string::npos) {
    VeloceMetrics::fallback_transfers_count.fetch_add(1, std::memory_order_relaxed);
    return real_cudaMemcpyAsync(dst, src, count, kind, stream);
}

uint32_t target_epoch = global_tx_sequence.fetch_add(1, std::memory_order_relaxed);
FlagPool::Slot& slot = g_gpu_flag_pools[active_gpu_ctx]->get(pool_idx);
*(slot.host_pinned_sig_ptr) = target_epoch;

// Intercept target stream context with hardware spin-kernel
typedef void (*ptx_kernel_t)(volatile uint32_t*, uint32_t);
void* kernel_ptr = (void*)dlsym(RTLD_DEFAULT, "veloce_ptx_sync_kernel");
if (kernel_ptr) {
    auto setup_kernel = (ptx_kernel_t)kernel_ptr;
    setup_kernel<<<1, 1, 0, stream>>>(slot.dev_flag_ptr, target_epoch);
}

int dynamic_relay_node = (active_gpu_ctx + 1) % 8;

Task memory_flight_task;
memory_flight_task.host_src = (uint8_t*)src;
memory_flight_task.dev_dst = (uint8_t*)dst;
memory_flight_task.total_bytes = count;
memory_flight_task.target_gpu = active_gpu_ctx;
memory_flight_task.assigned_relay_gpu = dynamic_relay_node;
memory_flight_task.flag_pool_idx = pool_idx;
memory_flight_task.epoch = target_epoch;
memory_flight_task.is_direct = false;

g_veloce_daemon->push_job(memory_flight_task);
return cudaSuccess;
}
}

```

5.2 Device-Side Control and Core Sync Kernels

The physical synchronization engine is implemented via inline PTX instructions inside the device synchronization module. It forces direct global memory transaction loads bypassing cache lines, keeping streaming processor utilization below 0.01%.

Listing 2: Device-side PTX synchronization kernel.

```

#include <cuda_runtime.h>
#include <stdint.h>

```

```

extern "C" __global__ void veloce_ptx_sync_kernel(volatile uint32_t* sync_flag, uint32_t expected_epoch) {
    if (threadIdx.x == 0 && blockIdx.x == 0) {
        uint32_t current_epoch = 0;
        uint32_t sleep_ns = 128;
        constexpr uint32_t MAX_SLEEP_NS = 2048;

        while (true) {
            // Read direct from HBM, bypassing standard L1 data cache architectures
            // PTX Instruction Target: ld.global.cg.volatile.u32
            asm volatile(
                "ld.global.cg.volatile.u32 %0, [%1];"
                : "=r"(current_epoch)
                : "l"(sync_flag)
                : "memory"
            );

            if (current_epoch >= expected_epoch) {
                break;
            }

            // Power-safe SM scheduling yield via assembly nanosleep
            asm volatile(
                "nanosleep.u32 %0;"
                :
                : "r"(sleep_ns)
            );

            sleep_ns = (sleep_ns * 2 > MAX_SLEEP_NS) ? MAX_SLEEP_NS : sleep_ns * 2;
        }
    }
}

```

6 Evaluation and Empirical Results

To evaluate VeloceIO under real production pressures, performance tests were performed on an 8× NVIDIA H20 cluster node featuring PCIe Gen 5 ×16 links and dual-socket AMD EPYC 9654 processors.

6.1 Raw H2D Streaming Scaling Analysis

We sweeping-test total achieved throughput against payload block capacity across an expanding number of active peer relay engines.

As detailed above, standard single-path execution achieves a maximum practical ceiling of 53 GB/s. With VeloceIO multiplexing enabled, throughput scales almost linearly before flattening at 245 GB/s when utilizing six or more relay nodes. This plateau matches the real physical limits of the dual-socket AMD xGMI3 inter-processor bridge links, approximately 256 GB/s bidirectional capacity limit.




Figure 2: Achieved H2D bandwidth as relay density increases.

6.2 Time-to-First-Token End-to-End Metrics

We evaluated user-facing serving metrics on vLLM engines running deep models, including DeepSeek-V3 and Qwen-72B, across varying context lengths.

Active Context Size	Native CUDA Baseline	VeloceIO Optimized Path	Effective Latency Speedup
32K Tokens Fetch	1.32 s	0.58 s	2.27× Faster Fetch
64K Tokens Fetch	2.62 s	1.10 s	2.38× Faster Fetch
128K Tokens Fetch	5.10 s	2.15 s	2.37× Faster Fetch
Model Cold Start (64GB)	2.10 s	0.88 s	2.38× Faster Load

Table 1: Time-to-First-Token and cold-start latency improvements.

6.3 Co-located Resource Contention Assertions

To measure multi-tenant isolation safety, large memory copies were executed over VeloceIO while simultaneously running compute-intensive matrix multiplication loops on adjacent relay GPUs.

Without the active yielding framework, unconstrained NVLink multi-path transfers cause memory controller contention, leading to a 14.8% execution slowdown on active processing layers. Activating the congestion-aware scheduler, Ψ_{yield} , successfully caps compute degradation to under 2.7%, verifying VeloceIO’s safety for shared enterprise environments.




Figure 3: Compute degradation under relay traffic with and without active yielding.

7 Discussion and Future Work

While VeloceIO achieves significant throughput gains inside single-chassis systems, extending its multipath paradigms to rack-scale, distributed environments presents an open research challenge. Future versions will explore mapping these techniques to disaggregated clusters using RoCEv2, InfiniBand RDMA, and CXL memory pools.

Additionally, we are tracking upcoming architectural shifts, such as NVIDIA’s Grace–Blackwell NVL72. These architectures connect the host CPU directly to the NVLink network, altering the traditional PCIe bottleneck. However, given the massive global deployment of standard PCIe-bound nodes, VeloceIO remains a critical optimization tool for modern enterprise hardware.

8 Conclusion

This paper introduced VeloceIO, a user-space, software-defined multipath memory access engine designed to eliminate the single-link PCIe bottleneck in multi-GPU nodes. By leveraging standard runtime interposition, VeloceIO parallelizes data movement across available peer PCIe links, utilizing high-speed NVLink backplanes for final assembly with zero application-level code modifications.

The empirical results demonstrate that VeloceIO increases effective host-to-device bandwidth from 53 GB/s to 245 GB/s. This translates to a $2.38\times$ reduction in Time-to-First-Token latencies for long-context inference and cuts serverless cold starts by over 75% while maintaining strict resource isolation for adjacent compute kernels. VeloceIO offers an efficient, transparent, and production-ready solution to accelerate large-scale AI serving infrastructure.